

EDGE—CLOUD // ORCHESTRATOR

Thesis Experimental Report

Research platform for intelligent Edge—Cloud AI orchestration
using NVIDIA Jetson and AWS

Scenario	High CPU Load
Description	Edge devices saturated (CPU 3–25%, mem 15–40%). Forces of
Seeds	2026, 42, 7, 1234, 9999
N (per seed)	100
Total decisions	2,000
Generated at	2026-07-08 21:22 UTC
Engines evaluated	Rule-Based · Decision Tree · Random Forest · Q-Learning

1 • Experimental Configuration

Decision Engine features:

- network_latency_ms
- connectivity
- cpu_available
- memory_available
- batch_size
- priority
- cost_budget_usd
- model_size_mb

Decision classes:

- edge
- cloud
- hybrid

Engines under study:

1. Rule-Based – deterministic heuristic baseline
2. Decision Tree – sklearn DecisionTreeClassifier (max_depth=8, min_samples_leaf=10)
3. Random Forest – sklearn RandomForestClassifier (n=80, max_depth=10)
4. Q-Learning – tabular contextual-bandit RL (5000 episodes, $\alpha=0.15$, $\gamma=0.9$, ϵ -greedy 0.9→0.05, 108 discretised states)

Execution measurement:

- latency: time.perf_counter (wall-clock)
- CPU/memory: psutil (real process metrics)
- cost: environment model derived from AWS EC2 inference pricing proxy

Statistical protocol:


- Same context vectors fed to all engines (paired-samples design)
- Multi-seed aggregation with Student-t 95% confidence intervals

2 · Comparative Results (mean ± 95% CI across seeds)


Policy	Latency p50 (ms)	Latency p95 (ms)	Cost total (\$)	Decision time (μs)	Reward	Agreement w/ Rule	Success rate
DecisionTree	193 ± 12	282 ± 7	0.754 ± 0.049	52.5 ± 12	-28.3 ± 0.93	0.936 ± 0.032	1 ± 0
QLearning	49.3 ± 0.65	60.7 ± 0.67	0.002 ± 0	22.2 ± 1.7	-3.41 ± 0.21	0 ± 0	1 ± 0
RandomForest	191 ± 10	282 ± 7	0.748 ± 0.049	1.3e+04 ± 6.2e+0	-27.8 ± 0.85	0.972 ± 0.022	1 ± 0
RuleBased	190 ± 10	282 ± 7	0.747 ± 0.05	3.51 ± 0.22	-27.8 ± 1	1 ± 0	1 ± 0

Seeds: [2026, 42, 7, 1234, 9999] · n per seed: 100 · scenario: high_cpu_load


3.1 · Latency distribution per policy




3.2 · Decision-time overhead




3.3 · Cumulative simulated cost



3.4 · Route distribution per policy



3.5 · Reward distribution



4 · Sample Explained Decisions

Example 1 – context:

network_latency_ms=67.9 · connectivity=1 · cpu_available=14.1 · memory_available=27.8 · batch_size=4 · priority=3 · cost_budget_usd=

- RuleBased → CLOUD (conf 0.90, 11.2µs) net<80ms; cpu low; mem low; model≥20MB
- DecisionTree → CLOUD (conf 1.00, 211.6µs) top-feature: network_latency_ms
- RandomForest → CLOUD (conf 0.81, 12372.8µs) top-feature: cpu_available
- QLearning → EDGE (conf 1.00, 81.6µs) state=9 · Q=[-2.49, -20.92, -11.34]

Example 2 – context:

network_latency_ms=132.2 · connectivity=1 · cpu_available=19.1 · memory_available=34.2 · batch_size=16 · priority=4 · cost_budget_usd=

- RuleBased → CLOUD (conf 0.85, 6.3µs) cpu low; mem low; batch≥8; prio≥4
- DecisionTree → CLOUD (conf 1.00, 192.1µs) top-feature: network_latency_ms
- RandomForest → CLOUD (conf 0.87, 11653.4µs) top-feature: cpu_available
- QLearning → EDGE (conf 1.00, 77.9µs) state=47 · Q=[-4.08, -40.91, -18.78]

Example 3 – context:

network_latency_ms=142.1 · connectivity=1 · cpu_available=5.4 · memory_available=35.4 · batch_size=8 · priority=2 · cost_budget_usd=

- RuleBased → CLOUD (conf 0.80, 6.7µs) cpu low; mem low; batch≥8; model≥20MB
- DecisionTree → CLOUD (conf 0.89, 207.1µs) top-feature: network_latency_ms
- RandomForest → CLOUD (conf 0.64, 11541.6µs) top-feature: cpu_available
- QLearning → EDGE (conf 1.00, 74.6µs) state=45 · Q=[-2.93, -29.32, -14.63]

5 · Methodology & Reproducibility

1. Scenario definition

Each experiment is executed under one of eight named scenarios (Mixed Traffic, Factory Normal, Network Failure, High CPU Load, Large Model, Cloud Congestion, Low Bandwidth, High Priority).

2. Reproducibility

Given (scenario, seed), the platform produces the exact same context sequence, so all four policies are evaluated on IDENTICAL inputs (paired-samples design). Multiple seeds are used to compute confidence intervals via the Student t-distribution.

3. Metrics collected

- Decision-time (μ s) – orchestrator overhead
- End-to-end latency (ms) – realised at the chosen route
- Cost (USD) – cumulative per policy
- Route distribution – % of edge/cloud/hybrid selections
- Success rate – feasibility of the chosen route in context
- Reward – combined objective
- Agreement w/ Rule – divergence from the deterministic baseline

4. Threats to validity

- Synthetic scenarios (mitigated by using diverse, semantically-named cases)
- Q-Learning is single-step (contextual bandit): consistent with the routing problem, but not full episodic RL
- Environment model estimates costs and latencies analytically

5. Artefact bundle

/tmp/benchmark/{benchmark_records.csv, benchmark_summary.json, benchmark_ci.json, chart_*.png, thesis_report.pdf}