


MASTER SUMMARY

Edge-Cloud AI Orchestration

Consolidated evaluation across all scenarios

Scenarios evaluated	8
Engines	DecisionTree · QLearning · RandomForest · RuleBased
Overall ranking (reward)	QLearning > RandomForest > RuleBased > DecisionTree
Generated at	2026-07-08 21:23 UTC

Scenario Wins by Policy (highest reward)



Global Average Metrics

Policy	Latency p50 (ms)	Cost total (\$)	Reward	Success rate	Agreement w/ Rule
DecisionTree	75.61	0.294	-11.363	100.0%	95.4%
QLearning	38.63	0.002	-0.658	100.0%	57.6%
RandomForest	71.95	0.276	-10.437	100.0%	93.3%
RuleBased	75.04	0.291	-11.160	100.0%	100.0%

Overall Ranking (by mean reward):

- 1. QLearning
- 2. RandomForest
- 3. RuleBased
- 4. DecisionTree

Per-Scenario Reward Ranking

Scenario	Winner	DecisionTree	QLearning	RandomForest	RuleBased
Mixed Traffic	QLearning	-11.82	-0.47	-10.75	-11.61
Factory Normal	QLearning	-1.14	0.68	0.03	-0.93
Network Failure	QLearning	-6.72	-0.21	-6.00	-6.72
High CPU Load	QLearning	-28.32	-3.41	-27.83	-27.77
Large Model	QLearning	-25.52	-0.77	-24.18	-25.24
Cloud Congestion	QLearning	-3.74	0.37	-3.73	-3.74
Low Bandwidth	QLearning	-0.66	0.01	-0.65	-0.65
High Priority	QLearning	-12.99	-1.47	-10.39	-12.62